


An i-vector based approach to accent conversion

Kenny W. Lino

M.Sc. Dissertation



Department of Intelligent Computer Systems
Faculty of Information and Communication Technology
University of Malta
2018

Supervisors:

Claudia Borg, Department of Artificial Intelligence, University of Malta
Andrea DeMarco, Institute of Space Sciences and Astronomy, University of Malta
Eva Navas, Department of Communications Engineering, University of the Basque
Country

Submitted in partial fulfilment of the requirements for the Degree of
European Master of Science in Human Language Science and Technology

M.Sc (HLST)
**FACULTY OF INFORMATION AND
COMMUNICATION TECHNOLOGY
UNIVERSITY OF MALTA**

Declaration

Plagiarism is defined as "the unacknowledged use, as one's own work, of work of another person, whether or not such work has been published" (Regulations Governing Conduct at Examinations, 1997, Regulation 1 (viii), University of Malta).

I, the undersigned, declare that the Master's dissertation submitted is my own work, except where acknowledged and referenced.

I understand that the penalties for making a false declaration may include, but are not limited to, loss of marks; cancellation of examination results; enforced suspension of studies; or expulsion from the degree programme.

Student Name: Kenny W. Lino

Course Code: CSA5310 HLST Dissertation

Title of work: An i-vector based approach to accent conversion

Signature of Student:

Date:

Abstract

With the emergence of the use of technology in language learning through tools like Rosetta Stone and Duolingo, learners have slowly been given more autonomy of their language learning projection. Although these tools have allowed learners to tailor their learning to their own liking, there is a gap between the available resources to assist those that would like to improve their pronunciation. Previous research in the intersection of language learning and speech technology has made efforts to develop pronunciation training systems to address this problem, but the systems themselves tend to have gaps due to the lack of appropriate support for the users, especially in appropriately identifying errors and providing sufficient feedback to help them correct their errors.

Some researchers have purported that alongside other forms of feedback such as a visual articulatory representation, a voice conversion system could serve as a potential feedback mechanism by helping learners understand what their voice could sound like given the appropriate changes. However, like pronunciation training systems, voice conversion systems also faced many limitations due to the complex interaction of various features which made them unrenderable as useful tools. With that said, recent advances in speech technology using methods such as i-vectors and deep neural networks have become increasingly successful in achieving better accuracy and quality in a variety of tasks, allowing for the potential to return and address these said gaps in performance for voice conversion.

This dissertation investigates these advancements in applying i-vectors and deep neural networks to develop an accent conversion system (a modified voice conversion system) that could potentially serve as a feedback mechanism as a part of a larger computer-based pronunciation training system. I compare this methodology against baselines using more traditional features and conversion processes following the work of Aryal and Gutierrez (2014) among their other works, and evaluate using the responses of participants in a perceptual study. I conclude with a discussion of the current work and highlight some potential directions for future direction.

Contents

Abstract	II
Contents	III
List of Figures	IV
List of Abbreviations	IV
1 Introduction	1
1.1 Research Questions	2
1.2 Thesis Overview	2
2 Background	3
2.1 Voice conversion	3
2.2 Accent conversion	6
2.3 Technical Background	6
2.3.1 Mel-frequency cepstrum coefficients	6
2.3.2 Gaussian mixture models	8
2.3.3 Neural networks	10
2.3.4 I-vectors	10
3 Literature Review	12
3.1 Theoretical and educational motivations	12
3.2 Spoken language technology for education	15
3.3 Voice conversion	18
3.4 Accent conversion	21
4 Design and methodology	27
4.1 Data	27
4.1.1 CMU ARCTIC corpus	27
4.1.2 L2-ARCTIC corpus	27
4.1.3 Experimental data set-up	28
4.2 Baselines: GMM-based voice and accent conversion	28
4.2.1 Tools	29
4.3 Experiment 2: I-vector based accent conversion	29

4.3.1	Tools	29
5	Evaluation and results	31
5.1	Evaluation	31

List of Figures

2.1	The training and conversion processes of a typical VC system.	5
2.2	The extraction of sequence 39-dimensional MFCC vectors from a waveform.	7
2.3	The Gaussian distribution with different means (μ) and variances (σ^2). .	8
2.5	Gaussian Mixture Model convergence using the Expectation-Maximization algorithm	9
2.4	An example of a GMM using male and female height. The likelihoods for each gender for someone 168cm (66in) tall is calculated using the percentage of men and women in the dataset from the vertical axis. The probabilities are given in the top right corner.	9
3.1	An example posteriorgram representation for the spoken phrase ‘basketball and baseball’.	24

List of Abbreviations

Keep this updated

CAPT	Computer Assisted Pronunciation Training
CP	Critical Period
GMM	Gaussian Mixture Model
L1/L2	First and second language
LSTM	Long-short term memory
MFCC	Mel-frequency cepstrum coefficient
TTS	Text-to-speech

Chapter 1

Introduction

Technology has continuously evolved to no bounds as witnessed by the current successes enjoyed by the use of neural networks and the power of current hardware, something perhaps predicted by Moore's Law who proclaimed that computing power would double once every 18 months (and then changed to 24 months) [cite here].

[Mention something about AI here]

We see the effects of neural networks throughout many subareas in computer science, including that of natural language processing. In fact, if we take a look at the number of publications involving neural networks, it has exponentially compounded annually [cite image here].

While technology has flourished and led to a number of new state-of-the-art systems such as improvements in commercial speech recognition and machine translation, it can be argued that these benefits have not reached and innovated other areas outside of research to the same extent. One such example that could benefit from modern innovations is education. Although there have been small trends here and there to create applications for educational use such as Duolingo for language learning [more examples?], in general it seems that education has not evolved at the same rate as technology. In particular, pronunciation has been a large standing challenge in language learning due to its complex nature. Unlike grammar and vocabulary, pronunciation can be challenging to both learn and teach due to the lack of clarity on how to teach it. Like grammar and vocabulary, pronunciation also involves a number of nuanced characteristics, including stress, rhythm, vowels, and consonants.

The variation in these features contribute to what many known as *accent*, or variations in pronunciation across speakers based on location, ethnicities, social classes, native languages, etc. Accents can be considered to be a part of dialects, where users of the same language may have variations beyond pronunciation, such as usage in vocabu-

lary or grammar. The line may often be blurred in everyday discussions and even in academic analyses as accent and dialect (as well as language) could be considered to be on a continuum, but for the sake of simplicity, I consider *accent* to be variations in pronunciation in this work.

1.1 Research Questions

In this thesis, I focus on investigating the following questions:

- How can we leverage recent advances in recent technologies (namely deep neural networks) to convert a speaker’s voice into sounding like it was said with another accent?
- What specific methodologies achieve the best similarity to the target accent and produce the most natural sounding audio?

1.2 Thesis Overview

The overview of the thesis is as follows:

In **Chapter 2**, I give a proper definition of voice conversion and accent conversion, and a high level overview of some technical details needed to better understand the current work.

In **Chapter 3**, I present the motivation for creating an accent conversion system by discussing previous findings in second language acquisition research especially in relation to speech. I then cover previous work in voice and accent conversion to frame the advances and shortcomings of previously developed systems.

In **Chapter 4**, the design and methodology of the experiments are presented alongside the appropriate tools utilized to conduct each one.

In **Chapter 5**, the results of the experiments previously described are presented along with some short discussion and conclusions drawn from the results.

In **Chapter 6**, the thesis is concluded with a reflection on the work presented along with some appropriate suggestions for future work.

Chapter 2

Background

Before delving into previous literature and their relevance to this work and the fields of NLP and language learning as a whole, I detail both voice conversion and accent conversion in order to help better distinguish them. I also go over some common speech technology concepts typically used in these systems at a high level in order to make the current work more accessible to those unfamiliar with the area. Further reference is also provided for those interested in the technical aspects and formalisms.

2.1 Voice conversion

To properly frame voice conversion, we take a look at Mohammadi and Kain (2017) who present a recent overview of the subfield. Following a definition set forth by the authors, voice conversion refers to the transformation of a speech signal of a *source speaker* to make it sound as if it were uttered by a *target speaker* in any chosen fashion with the utterance still being intact. Some of these changes can include changes in emotion, accent, or phonation (whispered/murmured speech). There have been a number of proposed uses for VC, including the transformation of speaker identity (perhaps for voice dubbing), personalized TTS systems, and against biometric voice authentication systems.

Voice conversion often involves a large number of processes, one of which includes deciding the appropriate type of data. To start, one must decide whether to have parallel or non-parallel speech data. Parallel speech data refers to speech data that has source and reference speakers that say the same utterance, so only the speaker-specific information is different, while non-parallel data would indicate datasets where the utterances are not the same, and thus entail further processes to create a target waveform. Even though parallel corpora are more desirable as it reduces the footprint necessary for con-

version, parallel corpora are often curated for specific purposes and are not available in most cases. Because of its simplicity, in some cases, researchers have tested making a psuedo-parallel corpus using acoustic clustering when working with non-parallel data (Lorenzo-Trueba et al. 2018; Sundermann et al. 2006).

Other aspects that need to be considered as discussed by Mohammadi and Kain (2017) include whether the data is *text-dependent* or *text-independent*. Text-dependent corpora indicate that the data has word or phonetic transcription, which can ease the alignment process during training, while systems using text-independent data would need to find similar speech segments, using a method like acoustic clustering before training. Finally, one minor aspect that is not considered often is the languages of the source speaker and target speaker. Although many systems tend to focus on voice conversion between two native speakers of the same language, systems that aim to convert between two speakers speaking in different languages would have to be wary of potential mapping issues between sounds. This is especially important to consider in terms of accent conversion, which will be discussed in the following section.

Aside from considering these aspects of the corpora, the type of features extracted from the waveforms heavily impact the quality of the conversions. In investigating the most salient features of speaker individuality, previous researchers have concluded that the average spectrum, formants, and average pitch level are the most relevant. Following these conclusions, most VC systems focus on converting these features, and often work at the frame-level (windows of $\sim 20\text{ms}$), with the assumption that the frame represents a stationary sound. From these frames, there are a number of common local features that are extracted to represent the signal. These include the spectral envelope, cepstrum, line spectral frequencies (LSF) and the aforementioned formants. On top of these local frame-based features, contextual features can be considered as well as the local features alone are often limited in what they can model. These contextual features can be as simple as adding delta and delta delta features, although methods such as something known as event-based encodings have been tested as well. With event-based encodings, a sequence of local features are separated into different event targets and transitions to model an utterance. However, this method faces the challenge of properly defining events within the sequence. Thus, although many algorithms and methods exist to model a signal, most systems focus on working with mel-frequency cepstrum coefficients (MFCCs) and deltas/double deltas, as they are very standard in most speech synthesis and recognition systems in general. The extraction process of MFCCs and deltas/double deltas are described in further detail in section 2.3.

After the chosen features are extracted, the features between the source speaker and target speaker have to be matched to prepare them for conversion. In parallel conversion, this means that each sound in an utterance has to be mapped between the speakers,

which can be done manually but more often is done using an algorithm such as dynamic time warping (DTW). Although this is usually an effective algorithm to find the best alignment, there can be issues in aligning the sounds as it assumes that the same phonemes of the speakers have similar features (Mohammadi and Kain 2017). This can be improved upon by adding phonetic transcription, or using methods such as forced alignment, but these methods may also have other limitations.

With non-parallel voice conversion, the alignment process becomes more complex as utterances from the source and target speakers have to be broken down into individual phonemes, and then the desired sounds must somehow be collected and synthesized to produce the converted speech. This can be done using methods like unit-selection text-to-speech (TTS), but this requires a large amount of annotated training data. Algorithms such as INCA can be used in addition to work without annotation by iteratively searching for the best frame pairs. Further information on the various alignment methods are detailed within Mohammadi and Kain (2017).

When the best frames between the source and target speakers are finally matched, a method has to be chosen to map the relationship between the frames. This has traditionally been done by using Gaussian Mixture Models, although neural networks have also become prevalent as well as they become ubiquitous throughout computational modeling. A detailed but accessible explanation of these algorithms and how they function is provided in section 2.3.

A visual representation that summarizes the voice conversion process can be seen in Figure 2.2, courtesy of Mohammadi and Kain (2017).

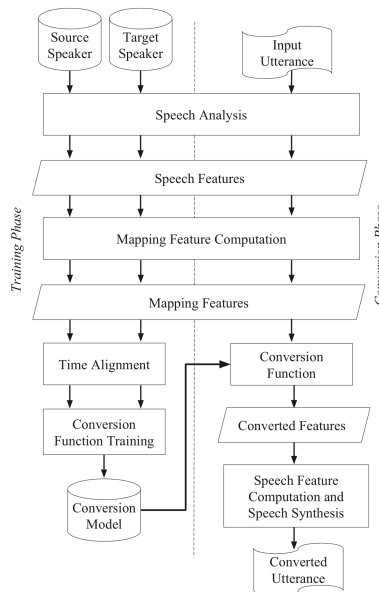


Figure 2.1: The training and conversion processes of a typical VC system.

2.2 Accent conversion

Like voice conversion, accent conversion is dedicated to convert the speech of a *source speaker* into sounding like a *target speaker*. However, accent conversion is specifically focused on morphing the *accent* of the speech signal, as opposed to sounding directly like the target speaker. Succinctly stated, “Accent conversion seeks to transform second language L2 utterances to appear as if produced with a native (L1) accent,” (Aryal and Gutierrez-Osuna 2014a). Because the confusion that can arise from using the terminology *source speaker* and *target speaker*, the *source speaker* is often referred to as the native or L1 speaker, while the *target speaker* is referred to as the non-native or L2 speaker. This seems somewhat counter-intuitive, but this allows for us to create a voice that retains the non-native speaker’s identity and the native speaker’s accent (Zhao, Sonsaat, Levis, et al. 2018).

Accent conversion poses a further challenge on top of (parallel) voice conversion as the audio of the source speaker and target speaker cannot simply be forced-aligned due to the fact that the voice quality and accent of the target speaker would remain (Aryal and Gutierrez-Osuna 2014b). This means that accent conversion may require more specialized alignment methods beyond standard frame-by-frame alignment that can help preserve the right speaker information while suppressing the other undesired information. This is further discussed in the examination of previous work in accent conversion in section 3.4.

2.3 Technical Background

2.3.1 Mel-frequency cepstrum coefficients

Following Jurafsky (2009), mel-frequency cepstrum coefficients (MFCCs) allow us to create vectorized representation of the acoustic information.

This is done by going over the speech signal using *windows*, where each window is assumed to contain a non-changing part of the signal. In other words, each window would roughly contain one phone– or speech sound. In order to retain all of the necessary information from each part of the signal, the windows often overlap.

After the signal is separated into different windows, the spectral information can be extracted using a special tool or formula known as the Discrete Fourier Transform. This allows us to find how much energy is in specific frequency bands.

From here the frequencies outputted by the Discrete Fourier Transform are converted

onto the *mel* scale, which is where the *mel* in Mel-frequency comes from. In short, the mel scale is used to represent human hearing, which is more sensitive to lower pitch sounds (under 1000hz) as compared to higher pitch sounds. Afterwards, the *cepstrum* is calculated in order to separate source information from filter information. From a high level, the source-filter theory says that all sounds come from the glottis (the area around our throat) and below, which contains information common to all speech sounds, such as the fundamental frequency (or pitch) of someone’s voice, as well as glottal pulse information. This is compared to the filter, which says that adjusting the vocal tract (e.g. moving the tongue and other articulators) define each individual sounds. By retaining just the filter information, we can model an individual phone. In terms of the given cepstral values, the first 12 cepstral values are taken as they neatly represent the filter information.

Although this information alone could be used to model a speech signal, additional information is often added to further better model each frame. Among this information is energy, which can help us further distinguish a sound, as vowels and sibilants (‘breathy’ sounds like /s/ or /f/) have more energy compared to stops (‘hard’ sounds like /k/ or /p/). On top of the 12 MFCC features and 1 energy feature, features known as deltas and double deltas are often added to represent the change in the speech signal frame to frame. Concretely, deltas can be used to model changes in formants or a change from stop closure to stop release. Double deltas are then added to represent the changes between deltas, which provide further precision in modeling an utterance. In total, this gives us 39 MFCC features from:

- 12 cepstral coefficients
- 12 delta cepstral coefficients
- 12 double delta cepstral coefficients
- 1 energy coefficient
- 1 delta energy coefficient
- 1 double delta energy coefficient

A visual representation of the whole MFCC extraction process can be seen in Figure 2.2, taken from Jurafsky (2009).




Figure 2.2: The extraction of sequence 39-dimensional MFCC vectors from a waveform.

2.3.2 Gaussian mixture models

A Gaussian mixture model is a type of probabilistic model that aims to represent normally distributed groups within a set. This is based on the idea of the normal, or *Gaussian* distribution, which can be seen in Figure 2.3.

[Create new graph from scratch or cite it.] The Gaussian distribution is characterized by two main features: the mean (the arithmetic average of the data) and the variance (the spread of the data from the mean). The Gaussian distribution is the most important distribution used in probabilistic modelling as it has been theorized that the average of independent random variables would look like a normal distribution (McGonagle et al. 2016).




Figure 2.3: The Gaussian distribution with different means (μ) and variances (σ^2).

Gaussian mixture models are based on the principle that if a unimodal (one ‘peak’) dataset can be fit with a Gaussian distribution, then a multimodal (multi ‘peak’) dataset is just a ‘mixture’ of smaller Gaussian distributions. A common example given to understand the Gaussian distribution and Gaussian mixture models often references height. It is often said that men are taller than women on average, with men being 178cm (5 foot 10 inches), and women being 165cm (5 foot 5 inches). If we used two separate Gaussians to model each gender, we could ‘mix’ them to model the likelihood of a certain data point (e.g. person) being a male or a female (McGonagle et al. 2016). For example, using a hypothetical example with the averages previously mentioned, we could see that the likelihood of a person that is 168cm is more likely to be a male than a female. This is demonstrated in Figure 2.4. [Create new graph from scratch or cite it.]




Figure 2.5: Gaussian Mixture Model convergence using the Expectation-Maximization algorithm




Figure 2.4: An example of a GMM using male and female height. The likelihoods for each gender for someone 168cm (66in) tall is calculated using the percentage of men and women in the dataset from the vertical axis. The probabilities are given in the top right corner.

However, as simple as this sounds the most advantageous point of the Gaussian mixture model is the fact that it is an *unsupervised* model that can be used when the subpopulations of the data are unknown. Thus, following the previous example of height, a Gaussian mixture model could be used to model the height of the two genders *without* knowing the gender of each datapoint.

Because it is an *unsupervised* model, it requires a special method to estimate the appropriate parameters. The most common method used for this is known as *expectation maximization*. This algorithm is used for maximum likelihood estimation. In other words, this algorithm tries to find the most appropriate group for each datapoint by calculating the probability of it being in a certain group and selecting the most likely one. This is done iteratively by initializing reasonable values, and then calculating the probability of membership in each cluster (the *expectation* step) and updating each clusters location, normalization and shape using the probabilities calculated (the *maximization* step) until the algorithms converge (VanderPlas 2016). A visual example of the convergence process taken from McGonagle et al. (2016) can be seen in Figure 2.5.

This model can be compared to the k -means clustering algorithm, as both can be used to cluster different subgroups. Like the k -means algorithm, GMMs also require us to specify a number of components, which usually indicate the number of subgroups we hope to cluster. However, k -means suffers from not using a probabilistic model to assign clusters, which means that data points can only be assigned to exactly one cluster. The cluster shape of k -means is also limited to only circles, which makes it inadequate to model data with different distributions. GMMs manage to address these issues by using the expectation-maximization algorithm to calculate the probabilities of cluster assignment and by allowing for different covariance types which permits for different cluster shapes beyond the circle. Aside from being useful as an unsupervised classification algorithm, GMMs can also be seen as a generative algorithm as it models the overall distribution of the data (McGonagle et al. 2016). This means that a GMM can be used to generate new data points following the distribution of the given data set.

In the case of speech, Gaussian mixture models are most often used to model individual sounds using MFCC feature vectors. Because MFCC feature vectors are multi-dimensional (~ 39 -dimensions), the Gaussians within the model are also multi-variate. However, the same principles described above still stand, and allow us to calculate the probability of a sound from a given frame.

Although GMMs are useful for modeling the distribution of sounds within a dataset and allow us to generate any observation, they are only capable of modeling speech as discrete MFCC vectors as opposed to a continuous sequence. Thus, GMMs are often utilized with Hidden Markov Models to remedy this issue. In short, Hidden Markov Models are models that consist of *states* and *transitions*. When utilized in speech recognition, each state represents a potential sound, while the transitions represents the probability or likelihood of the next state (e.g. sound). Gaussian mixture models can be used to represent a sound within a state [check from here; this also needs a diagram].

2.3.3 Neural networks

[Continue this part]

2.3.4 I-vectors

Identity vectors, or *i-vectors* are used as a method to model both the intra-domain and inter-domain variables of speech signals in the same low dimensional space. They are considered one of the more modern methods in speech recognition, compared to the more classic Gaussian mixture models and its modification, Gaussian Mixture Models

with universal background model (GMM-UBM), and joint factor analysis. In fact, i-vectors stem from a modification of the joint factor analysis method which is explained as the following.

[Continue this part]

Chapter 3

Literature Review

This section provides a brief overview of second language acquisition and education in order to frame the challenge of pronunciation and to motivate the potential usage of technology in language learning. I then examine some previous research in spoken language technology used in the domain of language education, including discussion on computer assisted pronunciation (CAPT) systems in order to shed light on where voice conversion and accent conversion could be applied, and then detail some important pivotal work done in the two areas.

3.1 Theoretical and educational motivations

Linguists have long debated over the possibility of whether second language (L2) learners (e.g. adult learners) could ever acquire a language to the extent of a native speaker. Some still cite ideas like the Critical Period (CP) Hypothesis and neuroplasticity which claims that learners cannot acquire language (at least as well as a native speaker) after a certain point in time due to the loss of plasticity in the brain (Lenneberg 1967; Scovel 1988). This theory has been particularly cited in reference to pronunciation, perhaps due to the obvious difficulty in overcoming the L1 negative transfer (e.g. effect of our native language) that many, if not all, language learners experience in speaking a new language.

Since the emergence of the CP hypothesis, many linguists have investigated the relationship between a number of variables such as age, motivation, and language use, that interact with the level of language acquisition. Piske et al. (2001) and Lengeris (2012) present an excellent review of different literature that investigates the interactions between these various variables and their effects on foreign accent. They discuss that although many L2 foreign accent studies do support the idea that the earlier a language

learner learns a language, the better their accent would be, there isn't strong enough indication to support the notion of a 'critical' period. They do concede that many studies do indicate that there is a linear correlation between age and foreign accent, but this only indicates a 'sensitive' period, not a 'critical' period, a distinction that some fail to acknowledge. That is to say, following advocates of the CP, the critical period should end roughly around 12 years old (Scovel 1988), or no later than 15 years old (Patkowski 1990), and beyond this point, there would be "a sharp drop-off in a learner's abilities" (Lengeris 2012), indicating that a learner could not acquire a native-like accent beyond this period.

Other researchers such as Long (1990) suggested that an L2 learner could speak accent-free if they learned the language before 6 years old but not after 12 years old. Although this notion also has been supported through a number of studies, there has also been counter-evidence found in other studies that found that there were learners younger than 6 who had detectable traces of a foreign accent. In other studies that examined learners of English who started beyond 12 years old, they also found evidence of learners with no detectable foreign accent. For example, in Flege et al. (1995), it was found that 6% of 120 native speakers of Italian who started learning English after the age of 12 years old had native-like pronunciation, and in Bongaerts et al. (1995), it was found that 5 out of 11 speakers were rated comparable to the native English control subjects. Thus Piske et al. (2001) conclude that while there is evidence that earlier learners can learn an L2 with less chance or degree of a foreign accent, this does not necessarily support the CP hypothesis or the idea that the loss of plasticity in the brain leads to an inability to acquire language.

Following their review of the correlation between other potential variables and the degree of foreign accent, Piske et al. (2001) take a further look at Bongaerts et al. (1995). Although it was found that 5 out of 11 speakers were perceived as *indistinguishable* from native English speakers, Piske et al. (2001) also point to a number of other factors in play that could have potentially allowed these speakers to be perceived as such. Perhaps the most controversial factor is the fact that these speakers were native Dutch speakers. Because Dutch and English have a similar phonetic inventory, it has been argued that this allows for easier acquisition of a native-like accent. In examination of their motivation, these learners were considered highly motivated as they felt the need to speak English without a Dutch accent due to their positions as university-level English teachers. On top of this, these learners also received intensive training in the perception and production of English sounds. Examining the interaction between these variables, it may be appropriate to conclude that the reason why 5 out of the 11 speakers were evaluated as indistinguishable from native speakers is due to the fact that they are a very niche population.

Regardless, this points to the fact that there is the possibility of language learners acquiring native-like pronunciation given the right configuration, even if the said configuration may be difficult to achieve. This is contrary to what is suggested by the CP hypothesis, which suggests that *all* learners are limited by a critical period. This is not to say that learners are not still deterred by a ‘sensitive period’, but this does highlight the potential that learners could be taught pronunciation should they have the right tools.

Aside from the issue of whether or not language learners could ever achieve native-like performance, another question that arises is whether or not there is even a *need* for learners to aim so high. In Munro and Derwing (1999), they discuss the interaction between foreign accent, comprehensibility and intelligibility and point out that the goal for many L2 learners is to communicate and not necessarily sound like a native speaker. Thus while there are unique groups of learners like those from Bongaerts et al. (1995), Munro and Derwing (1999) point out that most learners strive for effective communication. In order to observe the interaction between foreign accent, comprehensibility and intelligibility, they conduct a perceptual study on the performance of native Mandarin speakers. Following this study, they found that despite the fact that some speakers may have what some consider a ‘heavy accent’, this does not automatically mean that they are unintelligible. However, they do cite that some accents may cause longer processing times than others. When observing the interaction of variables such as phonemic errors and intonation with intelligibility and comprehensibility, they found that intonation was the most influential factor in comprehensibility, while phonemic errors affected intelligibility. This substantiates the concepts of comprehensibility and intelligibility themselves, as intelligibility is the degree a speaker is understood without involving interpretation (e.g. “What did they say?”), while comprehensibility is the degree a speaker is understood in terms of meaning (e.g. “What do they mean?”). Thus, they suggest that successful communication requires attention to both sounds and prosody for better comprehensibility and intelligibility.

While linguists make these discoveries and observations of L2 learning, it seems that it takes a lot of effort for them to trickle down to the foreign language classroom. In Darcy et al. (2012), they find through a small survey of 14 teachers that although teachers tend to find pronunciation to be ‘very important’, the majority do not teach it at all. When asked why they do not teach it, they cited reasons such as ‘time, a lack of training and the need for more guidance and institutional support’. Even though the number of teachers surveyed may be significantly small, this gives us a glimpse through the lens of what language teachers themselves experience in relation to pronunciation. We see that even though teachers would like to address it, this would require a restructuring in their curriculum and training– something that would undoubtedly take even more time

before students get more pronunciation attention. Compounded with the issue of time and the fact that not all learners need or want equal amount of pronunciation training, it may be unlikely to see such change in second language curriculum so soon.

This points to the potential solution of employing a technology-based system to improve pronunciation as learners could individually address their needs *outside* of the classroom.

3.2 Spoken language technology for education

Over the decades, as speech technology has slowly evolved and started to show its potential, many researchers have tried to test its limits by innovating a number of systems to address the challenge of pronunciation. Included in these systems are systems such as computer-assisted pronunciation training (CAPT) systems which attempt to tutor pronunciation through explicit teaching as well as more modern gamified techniques, which attempt to coerce language learners in to practicing pronunciation by making the process more engaging.

Among the two, CAPT systems have had more history due to the extra development and testing gamified techniques require. In fact, gamified systems can be considered a subclass of CAPT systems, as both require a fundamental setup in order to assist the language learner. In general, these systems utilize some form of automatic speech recognition (ASR) to record a speaker and compares their recordings (usually) with a native speaker gold standard. They also usually include a feedback mechanism with a combination of pitch contours, spectrograms or audio recordings to help the user adjust their pronunciation, with gamified systems including at least a point mechanism to motivate the user.

In order to understand the connection between language education and spoken language technology, we take a look at Neri et al. (2002) where we presented with a thorough overview between the two areas. Here, we see that aside from the classroom, there seems to be an issue in relating the findings of linguistics/language pedagogy with technology. Part of the reason, they suggest, stems from the fact that there are not ‘clear guidelines’ on how to adapt second language acquisition research and thus many CAPT systems ‘fail to meet sound pedagogical requirements’. They emphasize the need for the learners to have appropriate input, output, and feedback and exhibit how the systems available at the time were lacking. For example, they criticize some CAPT systems that were prevalent at the time including systems like *Pro-nunciation* and the *Tell Me More* series for utilizing feedback systems that give the users feedback in waveforms and spectrograms, which cannot be easily interpreted without training. Further, they

argue that although visual feedback has its merits, this kind of feedback suggests to the user that their utterance must look close to what is shown on the screen, which is not the case. An utterance can be pronounced perfectly fine, but look completely different from a spectrogram, and *especially* a waveform due to the number of features represented in each visualization, such as the intensity, which will indefinitely vary from user to user and the given exemplar. They conclude their article by making it a point to discuss recommendations for CAPT systems, by stating that they should integrate what has been found in research from second language acquisition, and to train pronunciation in a communicative manner to give context to the learners. They also point to the problematic area of feedback and advise that systems provide more easily interpretable feedback with both audio and visual information, and propose that systems give exercises that are ‘realistic, varied, and engaging’. Although this article was first published in 2002, it outlines a good fundamental structure that CAPT systems require— something that many systems still seem to be lacking.

In Eskenazi (2009), we are given a more technical review of CAPT systems with attention to the different CAPT system types and their limitations as well as some discussion on prosody detection.

The article explains that CAPT systems can be generally split into two main types: individual error detection and pronunciation assessment. As indicated, individual error detection systems are more focused on one particular aspect of the user’s speech, such as the phones or pitch, while pronunciation assessment systems are more designed to represent how a human would judge a non-native utterance.

Early individual error detection systems, including one of the author’s very own Eskenazi and Hansma (1998), started by using a variety of speech recognition techniques such as forced alignment or unconstrained speech recognition. They also worked with a variety of measures to detect the differences between the individual errors and gold standard. Some of these measures include hidden Markov model (HMM) based recognition scoring, a confidence score based system known as Goodness of Pronunciation (GOP), and Linear Discriminant Analysis (LDA). Each of these measures were found to somehow detect the users’ errors; however they suffer from issues like low precision or the need for a very homogeneous sample (e.g. Japanese speakers).

Although some of these early systems showed some signs of promise, they tended to over-simplify the issue of pronunciation training. Eskenazi (2009) makes a point of this by emphasizing the fact that improving non-native pronunciation is not simply a binary question of native vs. non-native. Instead the L1 of the system’s users must be considered as this in itself can greatly affect the evaluation. Eskenazi (2009) also points out that the level of language learning of the speakers can also impact the metrics

and success of the system as well, and thus an appropriate population must be selected carefully when building a CAPT system, especially when considering individual errors.

In examining previous CAPT systems, Eskenazi (2009) briefly discusses prosody correction, an often overlooked area as previously mentioned. Eskenazi (2009) points to some pivotal works that have used a variety of methods to address the issue, including systems that use Pitch Synchronous Overlap and Add (PSOLA) to resynthesize the prosody of users to help them hear what an appropriate utterance would sound like. This has been suggested to be a potentially effective feedback mechanism to employ in future systems, as it has been said that imitating one's own voice is the most effective (Felps, Bortfeld, et al. 2009). Among prosody correction systems, Eskenazi (2009) mentions two main types—those that use appropriate L2 phonological models and break prosody down into two levels—syllable-word and utterance-phrase, and systems that detect the 'liveliness' of a speaker. However, these systems require tuning of a variety of features including F0, power, duration or phonetic transcription, which makes it difficult to automatically create the necessary adjustments not just cross-linguistically, but across speakers as well. Eskenazi (2009) concludes by stating that although such limitations exist in these systems, the usage of ASR and other speech technologies has grown from such a sparse beginning, and that because the market appeal for such systems is large, they shall soon serve central roles in language education.

Aside from general discussion on CAPT systems, Chun et al. (2008) present a review of other technologies used in pronunciation training, with emphasis on feedback mechanisms used to train prosody. They discuss four main tools used in teaching prosody: 'visualization of pitch contours', 'multimodal tools', 'spectrographic displays' and 'vowel analysis programs'. Citing previous work, it appears that they suggest that the visualization of pitch contours is the most robust method of feedback for learners as it is the most intuitive and non-language specific. Aside from this however, they also discuss the potential of a multimedia approach used by Hardison (2005) that integrates both audio and video in a system called *Anvil*. Following this research, users of this system were able to generalize their training beyond a sentence level and were able to perform better at a discourse-level. This methodology encapsulates a good feedback mechanism as described by Neri et al. (2002) as it provides adequate feedback by being easily interpretable, stimulates both audio and visual channels and puts the language in context.

Chun et al. (2008) also discuss the two main methods of such prosody systems: one which utilizes isolated scripted sentences and the other utilizing imitation. They conclude that neither method is useful for generalizing to novel methods and suggest that the training should relate to the ultimate goal.

[This part needs polishing] Other information they provide in this article are prosody models used in previous studies. We see that some previous studies have focused on utilizing a variety of sentence types to teach prosody, contrasting *wh*-questions, echo questions, either-or questions and statements. Like the other articles, the works examined in Chun et al. (2008) gives us insight on potential ways to improve future CAPT systems, as we are shown exemplars of potential input and positive reinforcements in successful types of feedback for the user. They conclude that in order to create better pronunciation training systems, we should take advantage of recent technology.

In recent works related to gamified language teaching, (Tejedor-García et al. 2017) experiment with utilizing synthetic voices for corrective feedback in a pronunciation training tool. In their study, they use Google’s offline Android text-to-speech (TTS) system as feedback for B1 and B2 Spanish learners of English, and have them focus on the six most difficult pairs of vowels [insert IPA here?]. In order to train the users, the researchers first had them watch videos that describe the articulatory/perceptive features of the vowels, and had them listen to a number of minimal pairs produced by the TTS system in succession. Afterwards, they were asked to discriminate minimal pairs in a listening task and then asked to pronounce them.

From this study, they conclude that making use of common commercial TTS are beneficial for users and instructors alike as indicated by both the improvement in performance by the users and the feedback given by those involved in the experiment. This provides further support for works like (Felps, Bortfeld, et al. 2009), who demonstrated that accent converted speech also However, because the study was limited to individual words and only six pairs of vowels, further experimentation needs to be conducted in order to fully support their claim.

3.3 Voice conversion

There have been a number of efforts to design voice conversion systems using various methodologies. Much like the rest of the speech technology field, earlier voice conversion systems began with utilizing MFCCs and GMMs for conversion and slowly evolved towards utilizing more advanced features and adaptation techniques.

[Discuss Toda(2007) or something here that uses GMMs.]

One particular method that has been applied recently from other areas of speech technology is the usage of *i-vectors*. I-vectors are akin to word embeddings in text-based natural language processing tasks in the sense that i-vectors encapsulate any type of desired speech information in a vectorized fashion.

How are they different from MFCCs though, because MFCCs also vectorize speech information.

I-vectors have proven to be successful in a number of tasks, such as speaker verification, language identification, and native accent identification. Of particular interest is DeMarco and Cox (2013), who shows that [continue here]

In the instance of voice conversion, i-vectors are made of speaker super-vectors trained on GMMs and low dimensional features that represent an individual speaker's features (Wu et al. 2016). This is extracted per utterance and then averaged to form an i-vector that represents an individual speaker. In this way, a source speaker's i-vector can be approximated towards a target speaker's i-vector by a mapping function using neural networks, gaussian mixture models, or other appropriate algorithms.

The usage of i-vectors in voice conversion has been seen in works such as Wu et al. (2016) and Kinnunen et al. (2017). Following Kinnunen et al. (2017), the usage of i-vectors in voice conversion aligns perfectly with the task as it is highly similar to speaker verification; however instead of being a classification task (e.g. is this said speaker or not), voice conversion is a regression task. In Wu et al. (2016), they test and compare a variety of frameworks, such as a deep bi-directional long-short term memory neural (DBLSTM) network architecture, a DBLSTM combined with an average voice model, a DBLSTM combined with an average voice model retrained on some paralleled data, and another model which combined a DBLSTM, average voice model, and i-vectors. In order to evaluate these models, they provide both an objective evaluation using a measure known as mel-cepstral distortion (MCD) and a subjective evaluation rated on quality and similarity, which was decided by the votes of 20 listeners.

[expand on Wu et al. (2016) more]

Following the results of the subjective evaluation, they find that adding i-vectors to the DBLSTM and average voice model outperforms the DBLSTM and average voice model *without* i-vectors and that the DBLSTM and average voice model with i-vectors performs almost as well as the DBLSTM that was retrained. This underscores the need for i-vectors to capture speaker characteristics; without, the system highly underperforms and has poor quality and similarity.

As oppose to Wu et al. (2016) which utilizes an *eigenvoice* (or average voice), Kinnunen et al. (2017) supersedes Wu et al. (2016) by not requiring *any* parallel data.

Although most voice conversion systems have been successful at the general task, many of them suffer from low quality and/or low naturalness in their final outputs. For example, in listening to the audio of Wu et al. (2016)¹

[Maybe change footnotes to say listen to audio on the accompanying disk?]

¹ Visit <http://www.nwpu-aslp.org/vc/apsipa-jiewu-demo.pptx> to hear samples.

it is apparent that regardless of the low quality of the original source and target audios, the quality of the converted audio sounds muffled. This can be attributed to the various nuanced steps and features required to have high quality voice conversion.

For example, in a shared task dedicated to voice conversion, appropriately called *The Voice Conversion Challenge* where many leading research groups involved in speech technology around the world have submitted systems in attempts to tackle the issue. In the second iteration of the challenge Lorenzo-Trueba et al. (2018), the organizers proposed both a parallel and non-parallel version of the task, both of which were evaluated on natural and similarity using crowdsourcing.

The type of systems submitted to the 2018 edition of the task displays the current state of voice conversion and perhaps machine learning research in general as this year saw a huge increase in the number of systems using neural networks. However, it does not go without saying that there were indeed systems that used more traditional statistical methods, such as Gaussian Mixture Models (GMM) and one of its variations, differential GMM (DIFFGMM).

In order to evaluate the systems, a group of roughly 300 listeners were gathered to carry out a perceptual evaluation. The systems were evaluated on two main measures: naturalness, which was evaluated on a scale of 1 (completely unnatural) to 5 (completely natural); and similarity, which was evaluated using a same/different paradigm. Following the results, only one system, referred to as N10, was able to outperform the baseline in terms of naturalness (alongside the original source and target audios). When observing the performance of other systems in terms of similarity, we see about 5 out of 23 submitted systems outperforming the baseline. From this, we can conclude that it is easier to create a system with high similarity than high naturalness, which is consistent with other common systems.

In discussing the results of the N10 system, the authors credit the success of the system to the *hundreds of hours* of external speech data that was utilized to train a model to recognize content-related features, as well as manual fine-tuning. The creators of this system also made use of WaveNet, a novel high-fidelity vocoder and dozens of hours of clean English speech, which could also explain the success of their results. Thus, as previously discussed, we can conclude that creating a high-fidelity voice conversion requires not only appropriate fine-tuning of the data, but also a large amount of external data to support the system.

Thus, even though many systems were neural network based, only one neural network based system was able to outperform the GMM-based baseline, which could suggest that NN-based methods require proper fine-tuning of the hyperparameters.

Although we see limitations in the systems presented in The 2018 Voice Conversion Challenge, there have been other efforts to present high quality voice conversion systems in works such as and Nguyen et al. (2016) and Fang et al. (2018).

Fang et al. (2018) leverages a cycle-consistent adversarial network (CycleGAN) architecture, a variation of the recently trending generative adversarial network (GAN) architecture, which was originally used for unpaired image-to-image translation. For example, GANs have been shown to be able to convert images of zebras into horses, as well as winter into summer.

While not necessarily directly related to the standard idea of voice conversion, there have also been some incredible breakthroughs in systems set forth by research teams at Google Brain. One such system involves the Tacotron end-to-end system, which has been proposed to replace the current set-up of text-to-speech systems by reducing the amount of components (decoder, vocoder, etc. **[IS THIS TRUE?]**) into one piece. The researchers working on this system have recently revealed a impressive system that also takes advantage of deep neural networks to encode speaker characteristics into embeddings, which are then utilized to transfer style (Wang et al. 2018). They show how their system is capable of transferring a variety of emotions and accents, making the synthesized audio sound more human-like. Samples of these audios can be found **[put link here.]**

Even though these systems created by Google Brain are highly impressive, it is evident that the reason for the success of their systems is due to very fine-grained parameter tuning and the availability of large-scale, high quality data that many research institutions likely do not have access to or have funding for. For example, if we juxtaposed the audio from the Google Brain systems to the best performing system of the Voice Conversion Challenge 2018, we can still observe some disfluencies in the audios of the best system of the VCC 2018. Thus, it may be a long while before the general public has the ability to completely replicate such systems and before this work trickles in to the domain of accent conversion.

3.4 Accent conversion

Due to the specialized nature of accent conversion as compared to voice conversion, there are fewer articles and systems available for reference. In fact, most of the recent articles that are easily accessible on accent conversion were all published by the same group of researchers at Texas A&M University.

However, before the work of these researchers, works such as Yan et al. (2004) and

Huckvale and Yanagisawa (2007), explored manipulating various features in order to observe their relationship with a perceived accent. In Yan et al. (2004), they manipulate spectral features, intonation patterns and duration in order to observe their correlation across British, Australian and American accents.

Through an ABX perceptual test, they found that 75% of the synthesized utterances were evaluated as having the native accent, highlighting the potential for segmental accent conversion.

In Huckvale and Yanagisawa (2007)

[go into these works a little more]

Felps, Bortfeld, et al. (2009)

Felps and Gutierrez-Osuna (2010) ? [paper on evaluation]

Aryal, Felps, et al. (2010)

Felps, Geng, et al. (2012)

With that said, Aryal and Gutierrez-Osuna (2014b) and other works done by the group of researchers have made efforts to address the challenge. Throughout their research, they test a variety of methodologies, including accent conversion through voice morphing and articulatory synthesis. In the same work of Aryal and Gutierrez-Osuna (2014b), they propose a variation to standard forced alignment techniques used in voice conversion to pair frames based on acoustic similarity.

To achieve this, they first use dynamic time warping (DTW) to align parallel utterances from the L1 and L2 speakers in order to apply vocal tract length normalization to dampen the differences in pitch. They then extract sequences of 24 MFCCs per utterance, and cluster the MFCC vectors into 512 clusters using the k -means algorithm to easily find the most acoustically similar sound for each frame. The most acoustically similar frames are then calculated by finding the closest L2 cluster, and then selecting the most similar frame within the cluster. After the closest vectors are paired, they map the conversion using a GMM.

In order to evaluate their system, they had a group of 13 participants rate 12 utterances from the test set for their perceived accent (Which utterance was less accented?) and perceived speaker identity (Does utterance X sound more similar to A or B?). This system was compared to a standard voice conversion system that uses standard forced alignment and trained using GMMs. They found that comparing the AC system to the original L2 audio resulted in participants rating the converted audio as sounding less accented 86% of the time, while the VC system compared to the original L2 audio was rated at 91% of the time. However, when the converted audios from both systems were

compared, participants rated the AC system to be less accented compared to the VC system 59% of the time. It was also concluded that the AC system was more successful in retaining speaker identity, as the participants found the converted audio more similar to the L2 speaker 78% of the time. More interestingly, they found that the AC system was especially effective in converted utterances that are harder for the L2 speaker to pronounce. This was measured by examining the relationship between the number of phonemes that do not exist in the L2 language (in this case Spanish), and the number of listeners who judged the converted speech as sounding less accented. They found that there was a 0.86 correlation, indicating the robustness of the AC system. Thus, it appears that adjusting the alignment method to align acoustically similar sounds is a good start for accent conversion systems.

Aryal and Gutierrez-Osuna (2014a)

In Aryal and Gutierrez-Osuna (2015), we see a more novel method that looks beyond acoustic features to perform accent conversion. Citing the results of their previous work, they motivate the usage of articulatory gesture information in accent conversion reasoning that acoustic-based systems often struggle in the challenge of separating accent from speaker identity, which causes the accent converted audio to sound like a combination of the L1 speaker and L2 speaker. To do this, they propose a system that combines both the more standard acoustic information like aperiodicity, pitch and energy from the L1 speaker with articulatory information recording using an electromagnetic articulograph (EMA). Like many recent works, they test a DNN-based mapping function between the L1 and L2 data, which they compare to the previously popular GMM-based system.

In the evaluation of their system, they again use crowdsourced efforts to rate their system based on intelligibility, accentedness, and speaker identity. According to their sample size of 15 participants, they find that the DNN-based system was rated to have a 4.3 out of 7 in terms of intelligibility as compared to 3.84 out of 7 for the GMM-based system, proving that including articulatory gesture information and DNNs are more robust in this instance. The participants also rated the DNN-based system to be more native-like in 67% of cases as compared to the GMM-based system. With that said, the test set was only 15 sentences, which indicates that 10 out of 15 sentences were better with the DNN system; thus the test set used may be too small to draw hard conclusions. The most important conclusions drawn from their experiments was that of the voice identity assessment. In asking the participants to rate whether an MFCC compression and AC audio from the DNN and GMM-based systems, they found that the participants were fairly confident that the two audios were from the same person with both systems, with the DNN-based system outperforming the GMM-based system as before at a score of 4.3 out of 7 on average, and the GMM-based system at a score of 4.0. However, this is

difficult to compare to more common acoustic-only accent conversion systems, as this is not including in their evaluation. With that said, it may be possible to conclude that this would outperform acoustic-based systems, as they proposed this system to tackle flaws in their previous work.

Evidently, although including articulatory gesture information seems to improve the performance of accent conversion systems, as discussed in the closing remarks of their paper, recording articulatory gesture information can cost a great deal of money and time (Aryal and Gutierrez-Osuna 2015). Most publically (and privatized) speech corpora also do not include this type of information, meaning that experimenting with it in accent conversion at a broader scale is unfeasible. Thus, it is ambitious to accept adding articulatory information to accent conversion systems and further work needs to be done in order to scale standard audio-based speech corpora.

Departing from utilizing articulatory gesture information, Zhao, Sonsaat, Levis, et al. (2018) returns to a more simpler method similar to Aryal and Gutierrez-Osuna (2014b). However, instead of matching frames based on their *acoustic* similarity, they test matching frames based on their *phonetic* similarity. They do this by mapping the frames of each source and target speaker into something referred to as a *phonetic posteriorgram*. Following Hazen et al. (2009), a phonetic posteriorgram is ‘a time vs. class matrix that represents the posterior probability of each phonetic class for each time frame’. An example of a phonetic posteriorgram taken from Hazen et al. (2009) can be seen in Figure 3.1.




Figure 3.1: An example posteriorgram representation for the spoken phrase ‘basketball and baseball’.

The phonetic posteriorgrams are computed using a native English speaker-independent acoustic model and then the most similar source and target frames are matched by calculating something known as the Kullback-Leibler divergence (0 indicating similar or

same behavior, 1 indicating completely different) between the source and target posteriorgrams. After matching the frames, they train GMMs with 128-mixture components to model the distribution of the MCEPs to convert the speech. The performance of this proposed system is then compared to a standard voice conversion system using dynamic time warping to align the frames and the system described in Aryal and Gutierrez-Osuna (2014b).

Like the previous works of Aryal and Gutierrez-Osuna (2014b) and Aryal and Gutierrez-Osuna (2015), this work also approaches evaluation using a perceptual listening test to evaluate acoustic quality, speaker identity and accentedness. However, in this work, they evaluate over 50 test utterances using 30 participants, which better substantiates their results compared to the evaluation of 10-15 utterances by 10-15 participants in some of their older studies.

In terms of acoustic quality, they found that their proposed posteriorgram method received a score of 3.0 on a Mean Opinion Score scale of 1 to 5 (with 1 being ‘bad’ and 5 being ‘excellent’), as compared to a score of 2.6 using the method from Aryal and Gutierrez-Osuna (2014b) and 2.5 for standard voice conversion.

[continue]

Aside from the work conducted by the research group at Texas A&M University, it appears to be that there are not many, if any other researchers currently working in this subarea of accent conversion. This may be because voice conversion still leaves a lot to be desired itself, suggesting that most researchers may want to focus on perfecting standard voice conversion before attempting to tackle something more fine-grained. However, as research in voice conversion continues to expand, it also creates the potential to apply methodologies from voice conversion to accent conversion. Following the general methodologies of voice conversion, I hypothesize that it should be plausible to convert accents in a similar fashion and apply more recent innovations to propose state-of-the-art methods.

Chapter 4

Design and methodology

In this chapter, I introduce the dataset and tools utilized in the experiments, and detail the procedures carried out to conduct the accent conversion process.

Expand this section

4.1 Data

The main datasets utilized in the following experiments is the Carnegie Mellon University (CMU) ARCTIC corpus (Kominek and Black 2004) and the L2-ARCTIC corpus (Zhao, Sonsaat, Silpachai, et al. 2018), a non-native English counterpart to the CMU Arctic corpus.

4.1.1 CMU ARCTIC corpus

The CMU ARCTIC corpus was originally designed to have good phonetic (specifically diphone) coverage for speech synthesis.

4.1.2 L2-ARCTIC corpus

The L2-ARCTIC corpus currently contains 10 non-native speakers of Hindi, Korean, Mandarin, Spanish and Arabic, with a male and female speaker for each language. At the time of writing, the curators of the corpus are working to add an additional 10 speakers to the corpus by September 2018.

The original audio was sampled at 44.1 kHz, with each recording at roughly 3.7 seconds on average. In total, the duration of the corpus is 11.2 hours, with each speaker

recording an average of 67 minutes of audio, or the complete ARCTIC sentence prompt list of 1,132 utterances. However, some speakers did not read all of the sentences and some recordings were removed as they did not have appropriate quality.

In addition to the audio files, the corpus also includes word and phoneme-level transcriptions and manually annotated errors for a 150-sentence subset of the corpus, designed to be used in computer-assisted pronunciation training tools. Within the subset, there are 100 sentences uttered by all speakers, and 50 sentences that contain phonemes that are considered to be difficult based on a speaker’s L1. This also includes phone addition, phone substitution, and phone deletion annotations in ARPAbet format, as well as optional comments left by the annotators.

4.1.3 Experimental data set-up

In the context of the experiments described below, following Zhao, Sonsaat, Levis, et al. (2018) only 150 utterances are utilized, with the utterances from the L2-ARCTIC corpus downsampled to 16 kHz to match the quality of the CMU ARCTIC corpus. Although the sample size is very small compared to the actual size of the corpora, a small sample is chosen to acknowledge the often little amount of data that is available/acquirable in building these systems. The 150 utterances that are used are chosen at random, but are ensured to be 150 utterances that all speakers have recorded. Out of these 150 utterances, 100 are randomly chosen as training utterances while the other 50 are used test utterances.

The speakers utilized in the experiments are also limited to speakers BDL (male) and CLB (female) from the CMU ARCTIC database, who are the native reference speakers, while the non-native speakers chosen from the L2-ARCTIC corpus are the native Korean speakers (HKK, male; YDCK, female), Hindi speakers (RRBI, male; TNI, female), and Spanish speakers (EBVS, male; NJS, female). This is mostly similar to the datasets in Zhao, Sonsaat, Levis, et al. (2018), with the exception of the Korean female speaker (YDCK) in place of the male Korean speaker (YKWK), which is not included in the current release of the L2-ARCTIC corpus, and the replacement of the native male Arabic speaker (ABA) with the two native Spanish speakers.

4.2 Baselines: GMM-based voice and accent conversion

In order to understand more traditional mapping methods used in voice and accent conversion, I follow the methods described in [something about Toda (2007)] for voice conversion and reimplement the method described in Aryal and Gutierrez-Osuna

(2014b) which utilized frame matching based on acoustic similarity. Each serves as a baseline to be compared to the proposed experiment described in the next subsection.

[put link to subsection here.]

In reimplementing Aryal and Gutierrez-Osuna (2014b), certain features were removed—namely vocal length tract normalization and prosody modification. Although it is discussed that vocal tract length normalization allows for better frame matching, it was assumed that converting audio between speakers of the same gender would have less impact from differences in vocal tract length. Inspection of preliminary conversion audio without these features compared to conversion with these features as offered by Zhao, Sonsaat, Levis, et al. (2018) also suggested little to no impact.

4.2.1 Tools

In order to do GMM-based voice conversion, I utilize the `nnmnkwii`¹ Python package which provides fast and easy functions to train voice conversion systems conveniently based on [Toda(2007)]. Alongside this package, I also utilize a number of other packages that `nnmnkwii` is dependent on, including `pysptk`, a Python wrapper for the Speech Processing Toolkit, `pyworld`, a Python wrapper for `WORLD`, a well-known tool for high-quality speech analysis and acoustic feature extraction, `librosa`, another package for audio analysis, and the common `scikit-learn` machine learning package for GMM training.

For GMM-based accent conversion, I also make use of the `nnmnkwii` package and its dependencies, but in addition, I use a custom method written to find the most acoustically similar for each frame and convert the corresponding the frames instead of frames that are matched using dynamic time warping.

4.3 Experiment 2: I-vector based accent conversion

This experiment is motivated by the works presented in Wu et al. (2016) and Kinnunen et al. (2017) and should be considered the main experiment of this work. Due to their flexible nature, i-vectors are an appropriate method to capture the representation of an accent in a compact way.

¹Found at: <https://github.com/r9y9/nnmnkwii>

4.3.1 Tools

In order to do the i-vector based accent conversion, I first utilized the SIDEKIT Python toolkit to extract the MFCCs, create a UBM-GMM and extract the i-vectors to represent each accent.

Chapter 5

Evaluation and results

In this chapter, I detail how the previously described experiments are evaluated following previous work such as Zhao, Sonsaat, Levis, et al. (2018) and present their results.

5.1 Evaluation

Voice conversion and accent conversion systems can be evaluated using either: a) objective measures or b) subjective measures. With objective measures, evaluation can be difficult as it requires intricate formulas that do not necessarily extrapolate across datasets or even individual audios (Felps and Gutierrez-Osuna 2010). With subjective methods,

In both cases, accent conversion systems are often evaluated on three features: the acoustic quality, speaker identity, and accentedness of each converted audio.

In the case of my own experiments, I choose to evaluate using a perceptual study due to its reliability and because of the complexity of using objective measures. I adapt the method utilized in Zhao, Sonsaat, Levis, et al. (2018), which in turn was adapted from Aryal and Gutierrez-Osuna (2014b), another previous work from the same research group. This is done so the reimplementation of their baselines can be compared against the baselines in their work and so the i-vector method can be directly juxtaposed against both sets of baselines as well as the current best method of using phonetic posteriorgrams discussed in Zhao, Sonsaat, Levis, et al. (2018).

Specifically, I gather a group of [15-30?] listeners to listen to [15-150?] test samples, with [33%] taken from each system. They are asked to evaluate on the acoustic quality using a 5-point Mean Opinion Scoring system, with 1 representing ‘Bad’ and 5 representing ‘Excellent’ and the speaker identity on a voice similarity score rang-

ing from -7 representing ‘definitely different speakers’ to +7 representing ‘definitely same speaker’. They are then asked to evaluate the accentedness of the test samples using a preference test to compare the VC system vs. the i-vector system, the AC system vs. the i-vector system, and the original L2 audio vs. the i-vector system.

[How will this be compared to the posteriogram system?]

Bibliography

- Aryal, S., Felps, D., & Gutierrez-Osuna, R. (2010). Foreign Accent Conversion Through Voice Morphing, 5.
- Aryal, S. & Gutierrez-Osuna, R. (2014a). Accent conversion through cross-speaker articulatory synthesis. (pp. 7694–7698). IEEE. doi:10.1109/ICASSP.2014.6855097
- Aryal, S. & Gutierrez-Osuna, R. (2014b). Can voice conversion be used to reduce non-native accents? (pp. 7879–7883). IEEE. doi:10.1109/ICASSP.2014.6855134
- Aryal, S. & Gutierrez-Osuna, R. (2015). Articulatory-Based Conversion of Foreign Accents with Deep Neural Networks, 5.
- Bongaerts, T., Planken, B., & Schils, E. (1995). Can Late Starters Attain a Native Accent in a Foreign Language? A Test of the Critical Period Hypothesis.
- Chun, D. M., Hardison, D. M., & Pennington, M. C. (2008). Technologies for prosody in context: Past and future of L2 research and practice. (pp. 323–346).
- Darcy, I., Ewert, D., & Lidster, R. (2012). Bringing pronunciation instruction back into the classroom: An ESL teachers’ pronunciation ”toolbox”, 18.
- DeMarco, A. & Cox, S. J. (2013). Native Accent Classification via I-Vectors and Speaker Compensation Fusion, 5.
- Eskenazi, M. (2009). An overview of spoken language technology for education. *Speech Communication*, 51(10), 832–844. doi:10.1016/j.specom.2009.04.005
- Eskenazi, M. & Hansma, S. (1998). The Fluency Pronunciation Trainer, 6.
- Fang, F., Yamagishi, J., Echizen, I., & Lorenzo-Trueba, J. (2018). High-quality nonparallel voice conversion based on cycle-consistent adversarial network. *arXiv:1804.00425 [cs, eess, stat]*. arXiv: 1804.00425 [cs, eess, stat]
- Felps, D., Bortfeld, H., & Gutierrez-Osuna, R. (2009). Foreign accent conversion in computer assisted pronunciation training. *Speech Communication*, 51(10), 920–932. doi:10.1016/j.specom.2008.11.004
- Felps, D., Geng, C., & Gutierrez-Osuna, R. (2012). Foreign Accent Conversion Through Concatenative Synthesis in the Articulatory Domain. *IEEE Transactions on Audio, Speech, and Language Processing*, 20(8), 2301–2312. doi:10.1109/TASL.2012.2201474

- Felps, D. & Gutierrez-Osuna, R. (2010). Developing objective measures of foreign-accent conversion. *IEEE Transactions on Audio, Speech, and Language Processing*, 1030–1040.
- Flege, J. E., Munro, M. J., & MacKay, I. R. A. (1995). Factors affecting strength of perceived foreign accent in a second language. *The Journal of the Acoustical Society of America*, 97(5), 3125–3134. doi:10.1121/1.413041
- Hardison, D. M. (2005). Contextualized Computer-based L2 Prosody Training: Evaluating the Effects of Discourse Context and Video Input. *CALICO Journal*, 22(2), 16.
- Hazen, T. J., Shen, W., & White, C. (2009). Query-by-example spoken term detection using phonetic posteriorgram templates. In *In Proc. ASRU*.
- Huckvale, M. & Yanagisawa, K. (2007). Spoken Language Conversion with Accent Morphing, 7.
- Jurafsky, D. (2009). *Speech and language processing: An introduction to natural language processing, computational linguistics, and speech recognition* (2nd ed.). Prentice Hall series in artificial intelligence. Upper Saddle River, N.J.: Pearson Prentice Hall.
- Kinnunen, T., Juvela, L., Alku, P., & Yamagishi, J. (2017). Non-parallel voice conversion using i-vector PLDA: Towards unifying speaker verification and transformation. (pp. 5535–5539). IEEE. doi:10.1109/ICASSP.2017.7953215
- Kominek, J. & Black, A. W. (2004). THE CMU ARCTIC SPEECH DATABASES. *th ISCA Speech Synthesis Workshop*, 2.
- Lengeris, A. (2012). Prosody and Second Language Teaching: Lessons from L2 Speech Perception and Production Research. In J. Romero-Trillo (Ed.), *Pragmatics and Prosody in English Language Teaching* (Vol. 15, pp. 25–40). Dordrecht: Springer Netherlands. doi:10.1007/978-94-007-3883-6_3
- Lenneberg, E. H. (1967). The biological foundations of language. *Hospital Practice*, 2(12), 59–67.
- Long, M. H. (1990). Maturational constraints on language development. *Studies in second language acquisition*, 12(3), 251–285.
- Lorenzo-Trueba, J., Yamagishi, J., Toda, T., Saito, D., Villavicencio, F., Kinnunen, T., & Ling, Z. (2018). The Voice Conversion Challenge 2018: Promoting Development of Parallel and Nonparallel Methods. *arXiv:1804.04262 [cs, eess, stat]*. arXiv: 1804.04262 [cs, eess, stat]
- McGonagle, J., Pilling, G., Tembo, V., Chumbley, A., Ross, E., & Khim, J. (2016). Gaussian Mixture Model. <https://brilliant.org/wiki/gaussian-mixture-model/>.
- Mohammadi, S. H. & Kain, A. (2017). An overview of voice conversion systems. *Speech Communication*, 88, 65–82. doi:10.1016/j.specom.2017.01.008

- Munro, M. J. & Derwing, T. M. (1999). Foreign Accent, Comprehensibility, and Intelligibility in the Speech of Second Language Learners. *Language Learning*, 49, 285–310. doi:10.1111/0023-8333.49.s1.8
- Neri, A., Cucchiarini, C., Strik, H., & Boves, L. (2002). The Pedagogy-Technology Interface in Computer Assisted Pronunciation Training. *Computer Assisted Language Learning*, 15(5), 441–467. doi:10.1076/call.15.5.441.13473
- Nguyen, H. Q., Lee, S. W., Tian, X., Dong, M., & Chng, E. S. (2016). High quality voice conversion using prosodic and high-resolution spectral features. *Multimedia Tools and Applications*, 75(9), 5265–5285. doi:10.1007/s11042-015-3039-x. arXiv: 1512.01809
- Patkowski, M. S. (1990). Age and accent in a second language: A reply to James Emil Flege. *Applied linguistics*, 11(1), 73–89.
- Piske, T., MacKay, I. R., & Flege, J. E. (2001). Factors affecting degree of foreign accent in an L2: A review. *Journal of Phonetics*, 29(2), 191–215. doi:10.1006/jpho.2001.0134
- Scovel, T. (1988). *A time to speak: A psycholinguistic inquiry into the critical period for human speech*. Newbury House Publishers.
- Sundermann, D., Hoge, H., Bonafonte, A., Ney, H., Black, A., & Narayanan, S. (2006). Text-Independent Voice Conversion Based on Unit Selection. (Vol. 1, pp. I-81-I-84). IEEE. doi:10.1109/ICASSP.2006.1659962
- Tejedor-García, C., Escudero, D., González-Ferreras, C., Cámara-Arenas, E., & Cardeñoso-Payo, V. (2017). Evaluating the Efficiency of Synthetic Voice for Providing Corrective Feedback in a Pronunciation Training Tool Based on Minimal Pairs. (pp. 25–29). ISCA. doi:10.21437/SLaTE.2017-5
- VanderPlas, J. (2016). *Python data science handbook: Essential tools for working with data*. O'Reilly Media, Inc.
- Wang, Y., Stanton, D., Zhang, Y., Skerry-Ryan, R. J., Battenberg, E., Shor, J., ... Saurous, R. A. (2018). Style Tokens: Unsupervised Style Modeling, Control and Transfer in End-to-End Speech Synthesis. *arXiv:1803.09017 [cs, eess]*. arXiv: 1803.09017 [cs, eess]
- Wu, J., Wu, Z., & Xie, L. (2016). On the use of I-vectors and average voice model for voice conversion without parallel data. In *2016 Asia-Pacific Signal and Information Processing Association Annual Summit and Conference (APSIPA)* (pp. 1–6). doi:10.1109/APSIPA.2016.7820901
- Yan, Q., Vaseghi, S., Rentzos, D., & Ho, C.-H. (2004). Analysis by synthesis of acoustic correlates of British, Australian and American accents. In *2004 IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)* (Vol. 01, I-637-40 vol.1). doi:10.1109/ICASSP.2004.1326066

- Zhao, G., Sonsaat, S., Levis, J., Chukharev-Hudilainen, E., & Gutierrez-Osuna, R. (2018). Accent conversion using phonetic posteriorgrams, 5.
- Zhao, G., Sonsaat, S., Silpachai, A., Lucic, I., Chukharev-Hudilainen, E., Levis, J., & Gutierrez-Osuna, R. (2018). L2-ARCTIC: A Non-Native English Speech Corpus, 5.